


Text Retrieval and Search Engines

System Implementation: Fast Search

ChengXiang “Cheng” Zhai
Department of Computer Science
University of Illinois at Urbana-Champaign

System Implementation: Fast Search



How to Score Documents Quickly

General Form of Scoring Function

The diagram illustrates the general form of a scoring function $f(q, d)$ with several components and their relationships:

- Final score adjustment**: A box with two arrows pointing to f_a and $f_d(d)$.
- Weight aggregation**: A box with an arrow pointing to f_a .
- Weight a matched query term in d**: A box with two arrows pointing to $g(t_1, d, q)$ and $g(t_k, d, q)$.

The equation is:

$$f(q, d) = f_a(h(\underbrace{g(t_1, d, q), \dots, g(t_k, d, q)}_{\text{Weight a matched query term in d}}), \underbrace{f_d(d)}_{\text{Final score adjustment}}, \underbrace{f_q(q)}_{\text{Final score adjustment}})$$

A General Algorithm for Ranking Documents

$$f(q, d) = f_a(\mathbf{h}(g(t_1, d, q), \dots, g(t_k, d, q)), f_d(d), f_q(q))$$

- $f_d(d)$ and $f_q(q)$ are pre-computed
- Maintain a score accumulator for each \mathbf{d} to compute \mathbf{h}
- For each query term t_i
 - Fetch the inverted list $\{(d_1, f_1), \dots, (d_n, f_n)\}$
 - For each entry (d_j, f_j) , compute $g(t_i, d_j, q)$, and update score accumulator for doc d_j to incrementally compute \mathbf{h}
- Adjust the score to compute f_a , and sort

An Example: Ranking Based on TF Sum

$$f(d,q)=g(t_1,d,q)+\dots+ g(t_k,d,q)$$

$$\text{where } g(t_i,d,q) = c(t_i,d)$$

Query = “info security”

Info: (d1, 3), (d2, 4), (d3, 1), (d4, 5)

Security: (d2, 3), (d4,1), (d5, 3)

Accumulators:		d1	d2	d3	d4	d5
		0	0	0	0	0
info	(d1,3) =>	3	0	0	0	0
	(d2,4) =>	3	4	0	0	0
	(d3,1) =>	3	4	1	0	0
	(d4,5) =>	3	4	1	5	0
security	(d2,3) =>	3	7	1	5	0
	(d4,1) =>	3	7	1	6	0
	(d5,3) =>	3	7	1	6	3

Further Improving Efficiency

- Caching (e.g., query results, list of inverted index)
- Keep only the most promising accumulators
- Scaling up to the Web-scale? (need parallel processing)

Some Text Retrieval Toolkits

- Lucene: <http://lucene.apache.org/>
- Lemur/Indri: <http://www.lemurproject.org/>
- Terrier: <http://terrier.org/>
- MeTA: <http://meta-toolkit.github.io/meta/>
- More can be found at <http://timan.cs.uiuc.edu/resources>

Summary of System Implementation

- Inverted index and its construction
 - Preprocess data as much as we can
 - Compression when appropriate
- Fast search using inverted index
 - Exploit inverted index to accumulate scores for documents matching a query term
 - Exploit Zipf's law to avoid touching many documents not matching any query term
 - Can support a wide range of ranking algorithms
- Great potential for further scaling up using distributed file system, parallel processing, and caching

Additional Readings

- Ian H. Witten, Alistair Moffat, Timothy C. Bell: Managing Gigabytes: Compressing and Indexing Documents and Images, Second Edition. Morgan Kaufmann, 1999.
- Stefan Büttcher, Charles L. A. Clarke, Gordon V. Cormack: Information Retrieval - Implementing and Evaluating Search Engines. MIT Press, 2010.